

Deterministic evidence producer
(pure Python, offline)

**trusted,
never verified**



Evidence E
entities A_E · values V_E · anchor a^*

*TIER A: $\text{enum}(A_{E,f})$ bound into
the tool schema at call time*

untrusted

LLM narrator — any provider
answers via a submit_investigation tool



TIER B: deterministic verifier, plain code
(1) $x \in A_{E,f}$ (2) $|v - V_E| \leq 0.005$ (3) a^* addressed

*violation:
retry (≤ 2),
naming the
offending items*

*after the retry budget:
strip unsound, flag omission*

User
validated citations + verified claims,
then free text marked NOT verified

**unverified prose
still reaches
the user**

GUARANTEED: cited columns and claims $\leftrightarrow E$ (C1, C2)
NOT GUARANTEED: verdict, prose, anchor (C3, flagged); $E \leftrightarrow \text{world}$